

WebGPU Geant4-DNA: browser-native, kernel-fused Monte Carlo track-structure simulation for radiobiology, validated against Geant4 11.4.1

Ahmet Barış Günaydın
Independent researcher | webgpudna.com

June 8, 2026

Abstract

We present a WebGPU/WGSL re-implementation of Geant4-DNA track-structure physics and Karamitros Independent-Reaction-Time (IRT) radiolysis chemistry that runs entirely in a web browser, with no CUDA toolchain and no installation. The simulation uses a *kernel-fusion* architecture: one GPU thread per primary electron executes the full particle history inside a single fused compute dispatch, eliminating per-step dispatch overhead. We do not transpile Geant4; we re-implement its physics models from the Geant4-DNA source and the G4EMLOW 8.8 tabulated cross sections, and use a normally-compiled Geant4 11.4.1 as a validation oracle on a falsifiable, artifact-backed ladder. Cross sections match the reference to peak ratios of 0.975 to 1.000; the continuous-slowing-down range at 10 keV is $0.988\times$ Geant4. WebGPU track-structure tracking (Phases A+B) runs **455** $\times$ faster than Geant4 single-thread ($280\times$ vs. an 8-thread build) on an Apple M2 Pro, but we are explicit about scope: the Geant4 figure is whole-process wall-clock and the WGSL figure is dispatch-only. We measured the gap — Geant4 init plus DNA table-build is a negligible ~ 2 s (so event-loop-only the figure is $452\times$), the one real asymmetry being the ~ 6.8 GB of per-event ntuple I/O Geant4 writes and WGSL does not — and we note that the kernel-fusion architecture itself accounts for only $\sim 2\times$ of the speedup (the fused Phase A is 2% of the tracking time; Phase B is an un-fused wavefront). The honest like-for-like figure is the *end-to-end* **1.48** $\times$, where both pipelines do the same work and the IRT radiolysis chemistry on CPU dominates the wall-clock—we report all of these and are explicit about which is which. We also report the project’s honest limits: a 23% cascade-ionisation deficit traceable to an intentional excitation cross-section choice (the remaining SIGMA_EXC_SCALE divergence; the formerly-compounding RECOMB_BOOST fudge has been removed, recovering the count from 32%), and the finding that the *inter-track* chemistry contribution—which closes most of the residual 1 μ s gap against Geant4’s **chem6**—is *not* reachable browser-native, because cross-primary IRT is $O(N^2)$ for a point source and GPU step-by-step chemistry is bounded by the encounter radius to $\sim 50\,000$ time steps. These negative results are quantified and motivate a native runtime as the path forward. All experiments are reproducible via committed JSON artifacts.

1 Introduction

Geant4-DNA [9, 2, 10] is the CNRS/IN2P3-coordinated extension of Geant4 [1] for discrete, event-by-event Monte Carlo track-structure simulation in liquid water at radiobiologically relevant (eV–keV) energies, together with the subsequent radiolysis chemistry. It is the de-facto reference for mechanistic modelling of radiation-induced DNA damage.

GPU acceleration of this class of simulation is an active area: gMicroMC [14] and MPEXS-DNA [13] target track-structure on the GPU, while GGEMS [3] and GPUMCD [7] target macroscopic

dose. All of these are built on CUDA or OpenCL and require a workstation-class install. To our knowledge there is no open-source GPU Monte Carlo for radiation transport that runs *browser-native*—i.e. on the cross-vendor WebGPU API, with zero install, on any machine with a modern browser. That accessibility gap is the motivation here: a browser-native engine is trivially shareable, reproducible without a toolchain, and usable for teaching and rapid validation.

Our contributions are: (i) a complete WebGPU/WGSL re-implementation of the Geant4-DNA DNA_Opt2 physics chain plus Karamitros IRT chemistry; (ii) a kernel-fusion architecture that yields large speedups over single- and multi-threaded Geant4; (iii) a research-grade, falsifiable validation ladder against compiled Geant4 11.4.1, with every claim backed by a committed artifact; and (iv) an honest characterisation of where the browser-native approach reaches its limits, including a quantified negative result on the inter-track chemistry contribution.

2 Methods

2.1 Re-implementation, not transpilation

None of Geant4’s runtime executes on the GPU. Each physics model is re-implemented as flat WGSL arithmetic, using the Geant4-DNA C++ source as the algorithmic specification (sampling routines and formulae) and the G4EMLOW 8.8 data files as the cross-section tables, converted once into WGSL constant arrays. A normally-compiled Geant4 11.4.1 / G4EMLOW 8.8 install serves as the *validation oracle*: we run the `dnaphysics` example single-threaded, dump its ntuples, and compare the WGSL output against them—the same validate-against-compiled-Geant4 methodology used by gMicroMC, MPEXS-DNA and GGEMS.

2.2 Kernel-fusion architecture

The pipeline has three phases (Fig. 1): **Phase A (primaries)** is a single fused `@compute` dispatch with one thread per primary electron; that thread runs the entire history—Born ionisation, Emfietzoglou excitation, elastic and vibrational scattering—in a `for` loop, with no per-step dispatch. **Phase B (secondaries)** is a wavefront stepper (one dispatch per physics step advancing all live secondaries), which cannot be fused because the secondary count is unknown until Phase A completes. **Phase C (chemistry)** runs the Karamitros 9-reaction IRT in a Web Worker off the main thread (a GPU spatial-hash alternative exists but is less accurate at long times; see Sec. 4). Energy deposition into a 128^3 voxel grid uses fixed-point integer `atomicAdd` ($\times 100$ per eV), since WGSL atomics are integer-only.

2.3 Ported models and deliberate divergences

We port Born ionisation (5 shells), Champion tabulated / screened-Rutherford elastic, Sanche vibrational excitation (9 modes), and the Karamitros 9-reaction IRT chemistry (6 species, Smoluchowski totally-diffusion-controlled and Onsager-screened partially-diffusion-controlled types) [11, 12]. Two deliberate departures from the DNA_Opt2 defaults are documented (Table 1): we use the Emfietzoglou excitation model [5]—the model used by the Geant4-DNA constructor `G4EmDNAPhysics_option4`, which is unrelated to the similarly-named EM Standard list `G4EmStandardPhysics_option4`—rather than Born, because it yields the correct initial $G(H) = 0.33$; and one tuning scalar, `SIGMA_EXC_SCALE` = 0.5, scaling the excitation cross section. A second scalar formerly in the codebase, `RECOMB_BOOST`, had no physical basis; we measured it to be nearly inconsequential (E10r) and **removed it** (set to its neutral 1.0), so production, this paper, and the shipped demo all run a




Figure 1: Three-phase pipeline. Phase A is fully fused (one GPU thread per primary, full history in a loop); B and C cannot be.

Table 1: Deliberate divergences from Geant4-DNA DNA_Opt2, with the measured cost of each (experiment IDs reference § Numbers).

Item	DNA_Opt2	This work	Cost
Excitation	Born	Emfietzoglou	$\sigma_{\text{exc}} 2.55\times$ (E6b)
σ_{exc} scale	—	$\times 0.5$	net $1.27\times$ (E6c)
Recomb. boost	—	off ($\times 1.0$)	$\Delta 1.4$ pp RMS (E10r)
Chemistry pool	single	per-primary	$\Delta G(\text{H}_2)$ (E10f)
Dose reduction	fp64	fp32 atomics	MC-equivalent, not bit-exact

pipeline that is *parameter-free* in that knob. The removal was validated end-to-end: cascade ions *recover* from 0.677 to $0.766\times$ (E7d, it had been destroying autoionisation ions), the chemistry shifts by only 1.4 pp RMS while *improving* $\text{OH}/\text{e}_{\text{aq}}^-/\text{H}$ (E10r), and the SSB indirect/direct ratio *holds* in PARTRAC’s 2–3 band at 2.32 with no recalibration (E13d). The only remaining non-unity scalar, $\text{SIGMA_EXC_SCALE} = 0.5$, is a documented physics-data divergence (Emfietzoglou vs Born), not a tuning fudge.

2.4 WGSL constraints

WGSL forbids recursion, dynamic dispatch and dynamic allocation, and allows atomics only on u32. Consequently: all physics is inlined into `main()`; secondaries are deferred to a buffer + wavefront rather than recursed; dose uses fixed-point integers; and all memory is pre-allocated.

2.5 Validation protocol

The repository follows a six-level falsifiable ladder. Every experiment emits a JSON artifact carrying the git SHA, a named seed, WGSL shader hashes, a pass bar, and per-row observations; failed experiments are committed with `status:"fail"` and a diagnosis rather than rerun. All numbers below are drawn from this artifact set. The reference configuration is $N = 4096$ primaries at 10 keV on an Apple M2 Pro (apple/metal-3 adapter). A reproducibility caveat applies: fp32 atomic reductions are not order-deterministic across GPU vendors, so cross-hardware results are statistically equivalent within Monte Carlo noise but not bit-exact.

Table 2: Cross-section validation vs. G4EMLOW 8.8. Peak ratio is WebGPU/Geant4 at the worst single point; medians of the relative deviation are $\sim 10^{-3}$.

Channel	Model	Peak ratio	Exp.
Ionisation σ_{ion}	Born (5 shells)	0.9987	E1
Excitation σ_{exc}	Emfietzoglou (5 lev.)	0.9970	E2
Elastic σ_{el}	Champion / scr.-Ruth.	0.9751	E3
Vibrational σ_{vib}	Sanche (9 modes)	1.0000	E4




Figure 2: Energy-dependent CSDA closure under the joint fix: WebGPU/Geant4 range ratio at the eight ESTAR energies, pre vs. post fix (8/8 improved; artifacts E5b, E5d).

3 Results

3.1 Cross sections (Level 1)

All 9 cross-section channels match the G4EMLOW reference (Table 2). Sanche vibrational is bit-exact; the largest deviation is the Champion elastic peak ratio at 0.975.

3.2 Track structure (Level 2)

The continuous-slowing-down (CSDA) range at 10 keV is 2714.4 nm vs. Geant4's 2747.5 nm, i.e. $0.988 \times (3.59\sigma)$, with energy conservation at 100.0 % (E5). Across the eight ESTAR energies the deficit is energy-dependent and was closed monotonically by the joint fix: the 100 eV ratio rose from $0.587 \times$ to $0.736 \times$ and all eight energies improved, with the lift inversely proportional to the original deficit—the expected signature of the excitation cross-section correction (E5b, E5d; Fig. 2).

The mean free path is 5–11 % short of Geant4 across six energy bins (median 0.941, E6). In production (parameter-free) the full-cascade ionisation count is 389.9 per primary vs. Geant4's 509.2—a 23 % deficit (E7d), with an implicit W -value of ~ 25.6 eV vs. ICRU 31's 21.4 eV ([8], derived from the recovered count). Source archaeology found the deficit had two causes: the intentional excitation inflation channels energy away from ionisation (the residual 23 %), and RECOMB_BOOST= 2.0 had been *destroying* $\sim$ a third of the autoionisation ions via boosted geminate recombination. Setting it

Table 3: Track-structure metrics vs. Geant4 11.4.1 at 10 keV, $N = 4096$.

Metric	This work	Geant4	Ratio (exp.)
CSDA range	2714.4 nm	2747.5 nm	0.988 (E5)
Energy cons.	100.0 %	99.99 %	1.000 (E5)
Ions / primary (production)	389.9	509.2	0.766 (E7d)
W -value	26.89 eV	21.4 eV [†]	1.257 (E5c)
Sec. KE/primary	143.4	144.9	1.010 (E8)

[†] ICRU 31 low-LET liquid water reference.

to the physical 1.0 recovered them (E7d, $0.677 \rightarrow 0.766\times$) while keeping the chemistry parameter-free (E10r)—dissolving the former two-knob tradeoff rather than hiding it. The secondary-electron kinetic-energy spectrum at creation matches to 1.0 % overall, with 7/8 log-bins within 0.1–3.1 % (E8). Table 3 summarises.

3.3 Chemistry (Levels 3–4)

We report *parameter-free* chemistry yields: the empirical RECOMB_BOOST knob (Table 1), which has no first-principles basis, is set to its neutral value 1.0. Against Geant4’s `chem6` at matched 10 keV, the G -values at 1 μs (E10r) are then $G(\text{OH})$ $0.91\times$, $G(\text{e}_{\text{aq}}^-)$ $0.86\times$, $G(\text{H})$ $0.93\times$, $G(\text{H}_2)$ $0.74\times$, $G(\text{H}_2\text{O}_2)$ $0.69\times$ (Fig. 3), a five-species RMS deviation of 19.7 % at 1 μs .

Crucially, we measured that the knob is *not* load-bearing: enabling it ($\times 2.0$) changes the overall RMS deviation by only ~ 1.4 pp (to 18.3 %), and although it lifts $G(\text{H}_2)$ it does so by introducing a $1.10\times$ overshoot in $G(\text{H})$ while leaving $G(\text{H}_2\text{O}_2)$ unchanged—so we drop it and report the parameter-free values above (E10r). A first-principles replacement for its intended effect (super-excitation autoionisation, or an Onsager/Mozumder escape-probability treatment of geminate recombination) is left to future work (Sec. 5).

One result is energy-resolved and tuning-independent: $G(\text{e}_{\text{aq}}^-)$ is non-monotonic between 1 and 3 keV, a 12.5 % dip significant at 126σ by primary-bootstrap (Fig. 4, E10b), which `chem6` independently reproduces (E10d)—real track-end / spur-structure physics, not an artifact.

3.4 DNA damage (Level 5)

We report Level 5 as a *calibrated* damage model, not a validated prediction, and state the distinction plainly. On a 21×21 B-DNA fibre grid (3.89 Mbp target) the strand-break counts are the geometric backbone-reach counts multiplied by per-event break probabilities: the counts are the geometric backbone-reach counts multiplied by per-event break probabilities ($P_{\text{dir}} = 0.15$, $P_{\text{ind}} = 0.05$). In parameter-free production ($\text{SSB}_{\text{dir}} = 28$, $\text{SSB}_{\text{ind}} = 65$), the indirect/direct ratio is 2.32, within PARTRAC’s 2–3 band [6]. P_{ind} was *calibrated* (from the Geant4 default 0.4) to land in that band, so the ratio is a fit, not an independent test, and PARTRAC is itself a simulation—but it held in-band across the RECOMB_BOOST $\rightarrow$ 1.0 flip with no recalibration (E13d; it was 2.46 at 2.0), so it is at least robust to that change.

Against experiment-calibrated cellular yields (~ 35 DSB and ~ 1000 SSB per diploid cell per Gy for low-LET radiation [15]), the absolute per-Da yields are high by $223\times$ (SSB) and $796\times$ (DSB) *when normalised by the box-average dose* (0.243 Gy). This is entirely a point-source dose-normalisation artifact, which we quantify (E12-local): the validation dumps launch all primaries from the origin, so 98.1 % of the deposited energy — measured from the pre-chemistry radical-position dump — lands in the central 3 μm fibre-core cube. The dose the DNA actually receives




Figure 3: Radiolysis G -value time evolution—parameter-free WebGPU yields (RECOMB_BOOST = 1.0; lines) against Geant4 `chem6` endpoints at $1\ \mu\text{s}$ (open circles). Artifact E10r.

is therefore $\sim 238\ \text{Gy}$ (concentration factor $C \approx 981$), not $0.243\ \text{Gy}$. Re-normalised by this local dose, the absolute yields are $0.34\times$ (**direct SSB**), $0.82\times$ (**DSB**), and $1.28\times$ (**total SSB**) of experiment — within a factor of ~ 3 , not orders of magnitude. The one residual is the DSB/SSB *ratio* ($3.6\times$ high), which dose normalisation does not touch and which traces to the calibrated P_{ind} above. Level 5 thus predicts absolute DNA-damage yields to order-of-magnitude once the dose is scored where the DNA sits; the strand-coincidence ratio remains a calibration, not a prediction. (A cleaner follow-up, E12-bulk, would distribute the primary start positions across the box so that box-average and local dose coincide and the C correction is unnecessary.)

3.5 Performance (Level 6)

Phase A+B tracking is **455** $\times$ faster than Geant4 11.4.1 single-thread (635 ms vs. a 289.1 s median over 3 trials, E15b) and **280** $\times$ vs. an 8-thread build (178.0 s, E15c). Peak Phase A throughput is 538 947 primaries/s at $N = 16384$ (E15). We state two caveats so these numbers are not over-read. *First, on apple-to-apple scope*: the Geant4 wall-clock is whole-process while the WGS 635 ms is dispatch-only, but we measured the difference and it is small where it matters. A 16-primary init-probe runs in 3.2 s (E15-fair), so Geant4 process init plus DNA physics-table construction is only $\sim 2.1\ \text{s}$ (0.7 % of 289 s); the 289 s is $\sim 99\%$ genuine event-loop, and event-loop-only the speedup is $452\times$, statistically identical to $455\times$. The one material asymmetry is per-event ntuple I/O — Geant4 writes 6.8 GB to disk over the full run (measured: 1.65 MB/primary from 16- and 256-primary probes, near-linear; the row-wise merge is the likely cause of the poor $1.62\times$ MT-8 scaling), which the WGS dispatch does not; a no-ntuple Geant4 build would isolate how much wall-time this adds, so $452\times$ is a mild over-estimate of a compute-only comparison. (We note for transparency that an earlier draft mis-estimated the fixed overhead at $\sim 160\ \text{s}$ from the MT scaling alone; the direct measurement corrects this.) *Second, kernel fusion is not what delivers the large factor*: a within-WebGPU comparison of the fused kernel against a modelled per-step dispatch gives a **40** $\times$ factor *for Phase A* (E16), but Phase A is only 14.4 ms (2 %) of the 635 ms; Phase B is an un-




Figure 4: $G(e^-_{aq})$ at $1 \mu s$ vs. primary energy: the non-monotonic 1→3 keV dip (126σ , E10b), independently reproduced by `chem6` (E10d).

fused 2000-dispatch wavefront. Fusion’s contribution to the full tracking pipeline is therefore $\sim 2\times$ (1324 ms naive $\rightarrow$ 635 ms fused), with GPU data-parallelism over N primaries supplying the rest. The single honest like-for-like number is the *end-to-end* $1.48\times$, where both pipelines do the same work and the IRT chemistry on CPU dominates the wall-clock (194 s of 194.6 s); this motivates Sec. 4.

4 The inter-track chemistry limitation

The largest remaining gap vs. `chem6` is an *inter-track* effect: per-primary IRT partitioning misses cross-primary reactions that, when included, add $\Delta G(H_2) = +0.149$ at $1 \mu s$ —96 % of the residual implementation gap (E10f). We asked whether this is recoverable browser-native, and found it is not, by two independent routes.

Cross-primary IRT is $O(N^2)$ for a point source. The IRT reaction horizon is the diffusion length over the full $1 \mu s$ window, $R_{cut} = 1.45 + 2\sqrt{8D_{max}t_{max}} \approx 551$ nm, not the few-nm contact scale. At 10 keV the radicals occupy a $\sim 5 \mu m$ blob, so a 551 nm spatial hash spans most of it and prunes almost nothing; over 50 % of accepted cross-primary reactions involve pairs more than 100 nm apart (E10k). The single-pool scan is $\sim 10^{12}$ reaction-time samples (~ 70 hr in-browser) and the heap is ~ 2 GB.

GPU step-by-step (SBS) chemistry is encounter-radius-bounded. The GPU radiolysis codes use a step-by-step Brownian-motion model in which a reaction fires when two reactants fall within a Smoluchowski reaction radius (gMicroMC [14]; MPEXS-DNA adds a during-step reaction probability at small time steps [13]). We implemented this in WGSL and, as a refinement, a Green’s-function (Brownian-bridge) during-step reaction probability [4] from the step-by-step radiation-chemistry literature, and measured its convergence against the analytic Smoluchowski value for a single pair. The per-step cost is ~ 110 ms at $N = 4096$ (E10L), and the IRT-parity budget is ~ 1750 steps. But accuracy requires a time step that resolves the encounter radius σ , $\Delta t \lesssim \sigma^2/(2D) \approx 0.02$ ns, *independent of pair separation or density* (E10Q; Fig. 6)—so ~ 50 000 steps are needed regardless. Compaction lowers per-step cost but not step count (E10P). No SBS




Figure 5: Speedups (log scale): fused track-structure kernel vs. Geant4 single-thread (E15b) and 8-thread (E15c), the within-WebGPU kernel-fusion factor (E16), and the honest *end-to-end* figure (E15b), which is only 1.48× because the IRT chemistry runs on CPU.

schedule, full or hybrid, beats the IRT worker at the required accuracy on laptop WebGPU.

Both routes converge on the same conclusion: the inter-track physics requires a *native* runtime (e.g. Node/Deno + `wgpu-native`), where RAM is cheap (unblocking single-pool cross-primary IRT) or many cores parallelise the step count. The browser remains the validation and demonstration surface.

5 Discussion

Browser-native execution buys accessibility and reproducibility that the CUDA/OpenCL engines cannot: the simulation and its validation harness run on any modern GPU with no install, which is valuable for teaching, sharing, and independent re-checking of every artifact. The cost is the WebGPU feature envelope (integer-only atomics, no recursion) and—per Sec. 4—an inter-track chemistry ceiling that pushes the last increment of fidelity off-browser. The `RECOMB_BOOST` fudge has been removed (the pipeline is now parameter-free in it); the residual 23 % cascade-ionisation deficit is the remaining open item, with an identified, falsifiable next step (super-excitation autoionisation tuning, to let `SIGMA_EXC_SCALE` relax toward unity). Relative to the CUDA codes `gMicroMC` and `MPEXS-DNA`, our speedups come from kernel fusion on a single consumer GPU (`MPEXS-DNA`, for reference, reports up to 2900× over single-core Geant4-DNA on an NVIDIA TITAN V [13]); extending the same methodology to the EM Standard list `G4EmStandardPhysics_option4` for macroscopic photon dose is a natural, if substantially larger, direction.

6 Conclusion

A browser-native, kernel-fused WebGPU re-implementation reproduces Geant4-DNA track structure to within a fraction of a percent on cross sections and 0.988× on CSDA range, at 455×/280× the speed of single-/multi-threaded Geant4, with an honest, artifact-backed account of its limits—most notably the quantified result that the inter-track chemistry contribution is not browser-native




Figure 6: Kill criterion (E10Q): single-pair SBS/analytic reaction-probability ratio vs. Δt for four separations. The curves overlap, so the Δt requirement is set by the encounter radius σ , not by separation — far, dilute pairs need the same tiny steps as close ones.

and requires a native runtime.

Data and code availability

Source, the full validation ladder, and every JSON artifact are at <https://github.com/abgnydn/webgpu-dna>; a live build runs at <https://webgpudna.com>. A versioned Zenodo deposit accompanies release v0.4.0 (DOI to be inserted). Each experiment reruns via `npm run experiments - <id>`.

References

- [1] S. Agostinelli et al. Geant4—a simulation toolkit. *Nuclear Instruments and Methods in Physics Research A*, 506(3):250–303, 2003.
- [2] M. A. Bernal, M. C. Bordage, J. M. C. Brown, et al. Track structure modeling in liquid water: A review of the Geant4-DNA very low energy extension of the Geant4 Monte Carlo simulation toolkit. *Physica Medica*, 31(8):861–874, 2015.
- [3] J. Bert, H. Perez-Ponce, Z. El Bitar, S. Jan, Y. Boursier, D. Vintache, A. Bonissent, C. Morel, D. Brasse, and D. Visvikis. Geant4-based Monte Carlo simulations on GPU for medical applications. *Physics in Medicine & Biology*, 58(16):5593–5611, 2013.
- [4] P. Clifford, N. J. B. Green, M. J. Oldfield, M. J. Pilling, and S. M. Pimblott. Stochastic models of multi-species kinetics in radiation-induced spurs. *Journal of the Chemical Society, Faraday Transactions 1*, 82:2673–2689, 1986.

- [5] D. Emfietzoglou, F. A. Cucinotta, and H. Nikjoo. A complete dielectric response model for liquid water: A solution of the Bethe ridge problem. *Radiation Research*, 164(2):202–211, 2005.
- [6] W. Friedland, M. Dingfelder, P. Kunderát, and P. Jacob. Track structures, DNA targets and radiation effects in the biophysical Monte Carlo simulation code PARTRAC. *Mutation Research*, 711(1–2):28–40, 2011.
- [7] S. Hissoiny, B. Ozell, H. Bouchard, and P. Després. GPUMCD: A new GPU-oriented Monte Carlo dose calculation platform. *Medical Physics*, 38(1):754–764, 2011.
- [8] ICRU. Average energy required to produce an ion pair. Technical Report Report 31, International Commission on Radiation Units and Measurements, 1979.
- [9] S. Incerti, G. Baldacchino, M. Bernal, R. Capra, C. Champion, Z. Francis, P. Guèye, A. Mantero, B. Mascialino, P. Moretto, P. Nieminen, C. Villagrasa, and C. Zacharatou. The Geant4-DNA project. *International Journal of Modeling, Simulation, and Scientific Computing*, 1(2):157–178, 2010.
- [10] S. Incerti et al. Geant4-DNA example applications for track structure simulations in liquid water: A report from the Geant4-DNA project. *Medical Physics*, 45(8):e722–e739, 2018.
- [11] M. Karamitros et al. Modeling radiation chemistry in the Geant4 toolkit. *Progress in Nuclear Science and Technology*, 2:503–508, 2011.
- [12] M. Karamitros et al. Diffusion-controlled reactions modeling in Geant4-DNA. *Journal of Computational Physics*, 274:841–882, 2014.
- [13] S. Okada, K. Murakami, K. Amako, T. Sasaki, S. Incerti, et al. MPEXS-DNA, a new GPU-based Monte Carlo simulator for track structures and radiation chemistry at subcellular scale. *Medical Physics*, 46(3):1483–1500, 2019.
- [14] M.-Y. Tsai, Z. Tian, N. Qin, H. Yan, Y. Lai, S.-H. Hung, Y. Chi, and X. Jia. A new open-source GPU-based microscopic Monte Carlo simulation tool for the calculations of DNA damages caused by ionizing radiation—Part I: Core algorithm and validation. *Medical Physics*, 47(4), 2020.
- [15] J. F. Ward. DNA damage produced by ionizing radiation in mammalian cells: identities, mechanisms of formation, and reparability. *Progress in Nucleic Acid Research and Molecular Biology*, 35:95–125, 1988.